VOICE NOTES · FIDELITY · 1/10

A transcript can sound polished and preserve nothing.

Diarization done right — the part that keeps a note usable a year later.

10 slides · swipe

Anton Petnitsky · Python / agents / backend

Fluent text without speakers is a riddle.

Two people debating an architecture. One model. Thirty minutes of 'we should / I disagree / fine, but'. Without attribution, the transcript is a ghost story with good grammar.

Words without a speaker
are content without evidence.

Diarization in practice.

01 Resemblyzer embeddings per utterance

↓

02 AgglomerativeClustering on cosine distance

↓

03 Speaker 1 → me, remainder → 'Speaker N'

↓

04 Note stamps decisions against speaker

No pyannote, no gated models — pure local embeddings.

The thresholds that actually work.

```
$ clustering check

intra-voice distance  : 0.87 – 0.95

inter-voice distance  : 0.69 – 0.72

cut at 0.30 cosine distance
```

Calibrated on two TTS engines — not on the README.

Anton Petnitsky · Python / agents / backend

One voice is a feature, not a failure.

If the whole recording is one person, the diarizer should shut up. Single-cluster output means no fake 'Speaker 1/2' markers polluting a clean solo note.

Diarize only when there's something to separate.

Anton Petnitsky · Python / agents / backend

What breaks diarization silently.

01 Code-switching mid-sentence (Russian → English)

02 Café noise that embeds as a third voice

03 Very short utterances — resemblyzer needs meat

04 Two people with similar vocal prints

05 TTS-vs-real voices in the same meeting

Every 'the diarizer lied' is an acoustics story.

Anton Petnitsky · Python / agents / backend

What the LLM gets to see.

Tagged segments — speaker, timestamp, text. The model groups by speaker, orders decisions, and preserves attribution. It never gets to decide who said what.

Attribution first.
Summarization second.

The test of a good transcript.

Open it in six months. Can you tell who committed to what — or only what was said?

If only the second, the transcript was decor, not record.

Anton Petnitsky · Python / agents / backend

Minimum viable fidelity.

01 Every decision carries a speaker tag

02 Contested points are marked as contested

03 Numbers and names are exact — even when boring

04 Uncertainty is preserved, not normalized away

Polished is optional. Faithful is not.

Anton Petnitsky · Python / agents / backend

VOICE NOTES · KEEP THIS · 10/10

Save this for your next transcript review.

When your transcript reads like a novel, ask who said it. If nobody can answer, your transcription solved the easy 40% and skipped the load-bearing 60.

Save for later review

Anton Petnitsky · Python / agents / backend